



# How Goldman Sachs works with Java

Deanna Lin 林文

---

we  
**BUILD**

## **1 Engineering at Goldman Sachs**

## **2 Using Java and OpenJDK in Goldman Sachs**

## **3 Eclipse Collections**

- **What's Eclipse Collections?**
- **Benefit of Eclipse Collections compared to Streams and JCF**
- **Learning Eclipse Collections by solving problems**

# Engineering at Goldman Sachs



# Goldman Sachs

---

**The Goldman Sachs Group, Inc.** is a leading global investment banking, securities and investment management firm that provides a wide range of financial services to a substantial and diversified client base that includes corporations, financial institutions, governments and individuals. Founded in 1869, the firm is headquartered in New York and maintains offices in all major financial centers around the world.

## **Goldman Sachs Technology Division**

We build solutions to some of the most complex problems, creating transformational technologies and developing systems that drive business and financial markets worldwide.

# Engineering at Goldman Sachs

## Leading latest open technology (global initiatives)

### Open Compute Project

- A project to build effective computation infrastructure at the lowest possible cost, resulting with optimized data center operation.

### Open Container Initiative

- A project to define format and runtime of container technology. More than 30 organizations including Docker are participating in the project.

### Java Community Process – Executive Committee

- JCP Executive Committee is responsible of important specification process for Java platform. Goldman Sachs was elected to the committee in 2011 for the first time and continue to be since then.

# Engineering at Goldman Sachs

**BUILD SOLUTIONS  
TO COMPLEX PROBLEMS**

LEARN MORE

we  
**BUILD**

# Using Java in Goldman Sachs

# History of Goldman Sachs and Java

1998

- Began internal Java evaluation

2000

- Java became one of the strategic platforms

2000-

- Continuously working closely with Sun/Oracle to improve Java

2004

- Ready for Lambdas in GS Collections Framework predecessor

2012

- Elected into JCP Executive Committee
- GS Collections Framework Open-sourced

2013

- Signed OCA
- Started contributing to OpenJDK

2015

- Migrated GS Collections to Eclipse Foundation



- JCP Executive Committee member
- Part of the Expert Groups for the following JSRs:
  - JSR 107: JCACHE – Java Temporary Caching API
  - JSR 335: Lambda Expressions for the Java Programming Language
  - JSR 351: Java Identity API
- Over 3000 Java Developers
- Maintain >130 million lines of Java code
- On average, ~125,000 Java processes running at any given hour every day

# Cutting Edge Java in GS

## ***We do push the boundary of Java technology***

- Very large heap
  - Some applications use as much as 190G heap
  - <https://bugs.openjdk.java.net/browse/JDK-6806226>
  - Integer overflow issues with GrowableArray caused issue with Mark Stack
- JVM stressed under high volume conditions or unusual usage pattern
  - Code cache capacity

## ***Open Source internally developed framework***

- Eclipse Collections: Collection framework that is richer and more fluent than Stream API

# Why We Use OpenJDK

- Source access is important!
  - Timely resolution for production issues
  - Identify potential area for improvement
  - Good for security
- Transparency
  - How is the code being developed?
  - Bug database
- Contribute patches for bugs and enhancements that benefit us and the wider Java community

# Typical Usage of OpenJDK

- **Troubleshooting**
  - Resolve production issues
  - Useful during development
- **Research new features**
  - Identify optimization opportunities in the JDK
  - How they can be applied and beneficial to applications

# Troubleshooting

***Having access to JDK source code is important***

- Timely Diagnostics
  - Crash dump analysis
  - Debug builds
- Identify workarounds and potential fixes
  - Create patched builds for developers to verify fixes
    - Including patches not yet available to GA builds

- Understanding new features in hotspot
  - How does it work?
  - Does it benefit our applications?
  - If it does, how could applications exploit it?
- Identify performance bottleneck
  - Potential workaround and features to improve application performance

# Case Study – Enhancement for OpenJDK

## Improve `Array.sort()` logic (DualPivotQuicksort) for primitive collections

### How

Changed the algorithm to

- reduce number of “run” (a “run” is how long you go through the array with it being in order)
- Stop sorting in the case of descending numbers

### Test case

- Additional functional unit tests
- Performance benchmark using JMH

### Patch contribution

- Discussed patch at <http://openjdk.5641.n7.nabble.com/Patch-to-improve-primitives-Array-sort-td224296.html>
- Patch was accepted in JDK 9 development branch
- In [JDK-8080945](#), the enhancement is included in June 2015

# Eclipse Collections

Open for contributions!

<https://github.com/eclipse/eclipse-collections/blob/master/CONTRIBUTING.md>



# What is Eclipse Collections?

## GS Collections

- Open source Java collections framework developed in Goldman Sachs
- Inspired by Smalltalk Collections Framework
- In development since **2004**
- Hosted on GitHub w/ Apache 2.0 License in **January 2012**

## Eclipse Collections 7.0

- Project @ Eclipse Foundation <https://www.eclipse.org/collections/>
- Feature set is the same as GS Collections 7.0
- Public released on Jan 20<sup>th</sup>!
- Packages renamed **com.gs** -> **org.eclipse**



# JavaOne2015 – Eclipse Collections by Example



# Eclipse Collections: Rich set of features

## Eclipse Collections 7.0

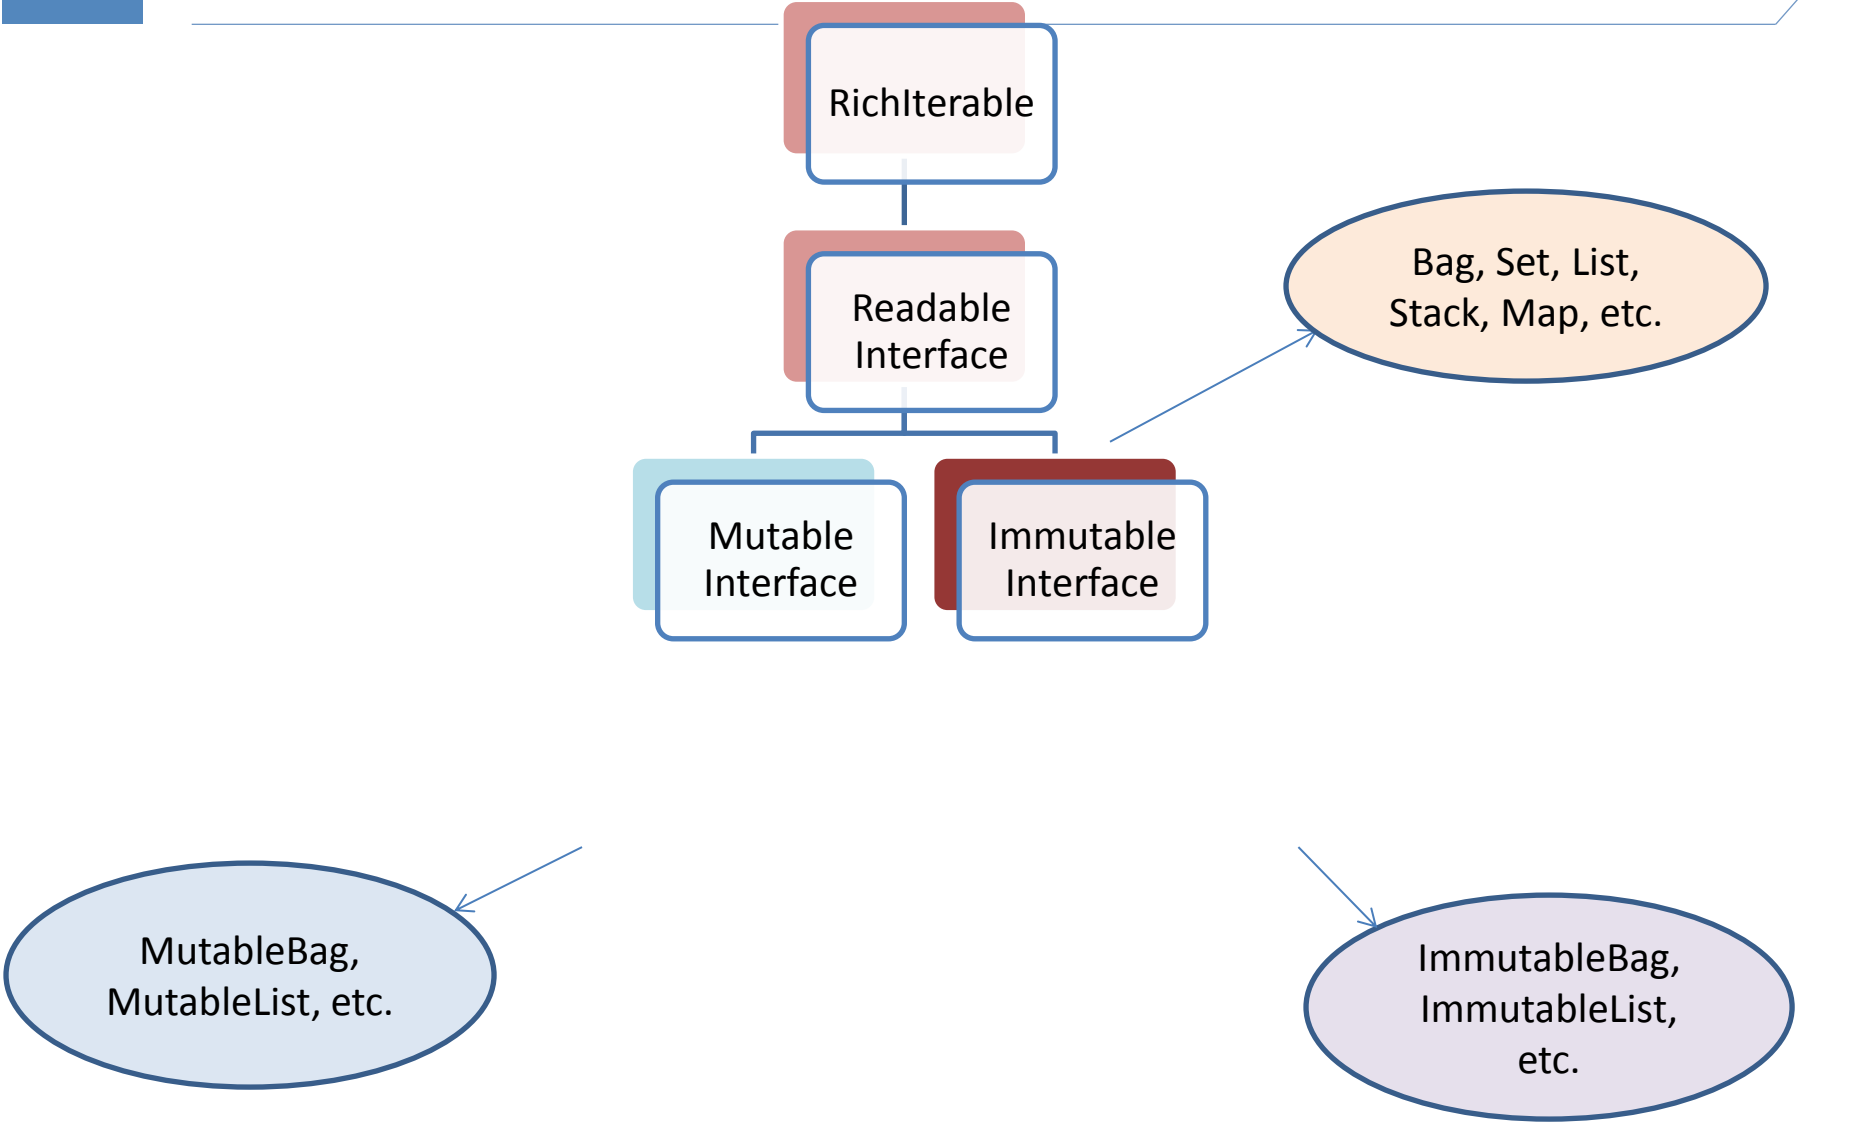
### Java 8 Stream API

- Functional APIs
- Lazy only
- Single use
- Serial & Parallel
- Primitive streams (3 types)
- Extensible Collectors

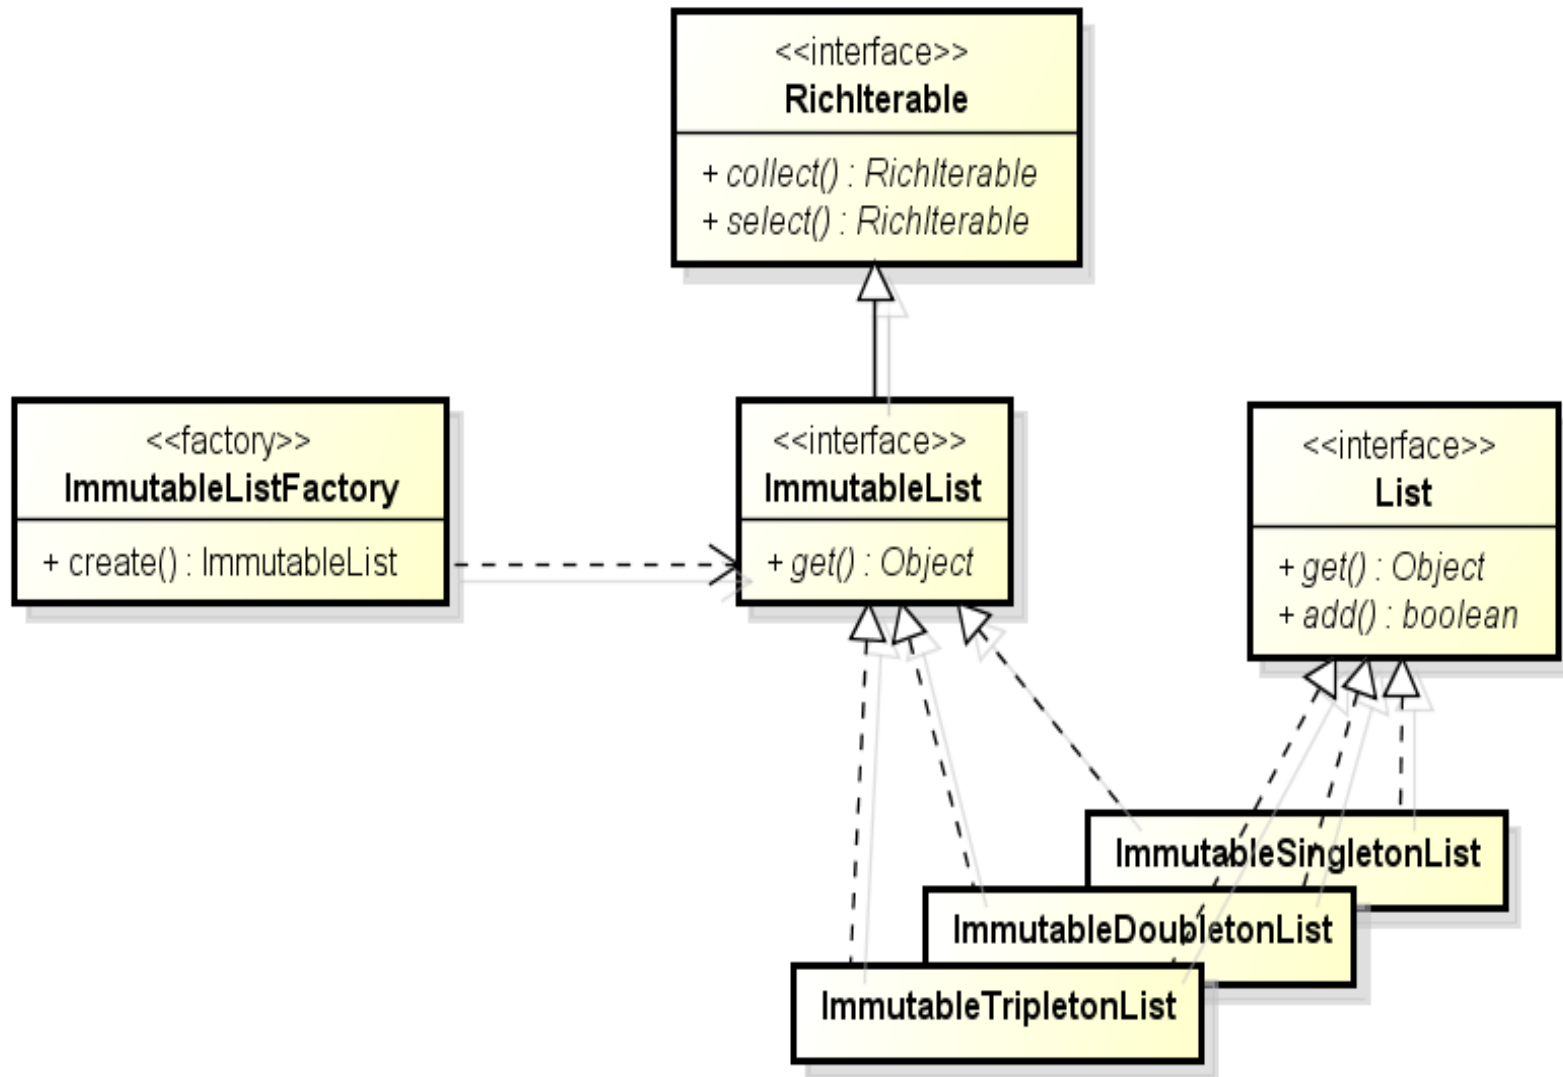
- Eager & Lazy, Serial & Parallel
- Memory efficient containers
- Primitive containers (all 8)
- Immutable containers
- More container types
- More iteration patterns
- “With” method patterns
- “target” method patterns
- Covariant return types
- Java 5+ compatible



# Design Concept



# Inheritance diagram - Immutability



# Collection Framework Comparisons

Features	Eclipse Collections	Java 8	Guava	Trove	Scala
Rich API	✓	✓	✓		✓
Interfaces	Readable, Mutable, Immutable, FixedSize, Lazy	Mutable, Stream	Mutable, Fluent	Mutable	Readable, Mutable, Immutable, Lazy
Optimized Set & Map	✓ (+Bag)			✓	
Immutable Collections	✓		✓		✓
Primitive Collections	✓ (+Bag, +Immutable)			✓	
Multimaps	✓ (+Bag, +SortedBag)		✓ (+Linked)		(Multimap trait)
Bags (Multisets)	✓		✓		
BiMaps	✓		✓		
Iteration Styles	Eager/Lazy, Serial/Parallel	Lazy, Serial/Parallel	Lazy, Serial	Eager, Serial	Eager/Lazy, Serial/Parallel (Lazy Only)

# Benefit of Eclipse Collections compared to Streams and JCF

# Write Concise Code

## Stream API

```
boolean result = this.people.stream().anyMatch(person -> person.hasPet(PetType.CAT));  
  
long result = this.people.stream().filter(person -> person.hasPet(PetType.CAT)).count();  
  
List<Person> peopleWithCats = this.people.stream().filter(person -> person.hasPet(PetType.CAT))  
    .collect(Collectors.toList());
```

## Eclipse Collections

```
boolean result = this.people.anySatisfy(person -> person.hasPet(PetType.CAT));  
  
int result = this.people.count(person -> person.hasPet(PetType.CAT));  
  
MutableList<Person> peopleWithCats = this.people.select(person -> person.hasPet(PetType.CAT));
```



# Leverage method references

## Stream API

```
boolean result = this.people.stream().anyMatch(person -> person.hasPet(PetType.CAT));  
  
long result = this.people.stream().filter(person -> person.hasPet(PetType.CAT)).count();  
  
List<Person> peopleWithCats = this.people.stream().filter(person -> person.hasPet(PetType.CAT))  
    .collect(Collectors.toList());
```

## Eclipse Collections

```
boolean result = this.people.anySatisfyWith(Person::hasPet, PetType.CAT);  
  
int result = this.people.countWith(Person::hasPet, PetType.CAT);  
  
MutableList<Person> peopleWithCats = this.people.selectWith(Person::hasPet, PetType.CAT);
```

# Code with Better Readability

## Stream API

```
List<Person> peopleWithCats = this.people.stream().filter(person -> person.hasPet(PetType.CAT))
    .collect(Collectors.toList());

List<Person> peopleWithoutCats = this.people.stream().filter(person -> !person.hasPet(PetType.CAT))
    .collect(Collectors.toList());

Map<Boolean, List<Person>> peopleWithcatsAndNoCats = this.people.stream().collect(
    Collectors.partitioningBy(person -> person.hasPet(PetType.CAT)));
peopleWithcatsAndNoCats.get(true);
peopleWithcatsAndNoCats.get(false);
```

## Eclipse Collections

```
MutableList<Person> peopleWithCats = this.people.select(person -> person.hasPet(PetType.CAT));

MutableList<Person> peopleWithoutCats = this.people.reject(person -> person.hasPet(PetType.CAT));

PartitionMutableList<Person> peopleWithcatsAndNoCats = this.people.partition(
    person -> person.hasPet(PetType.CAT));
peopleWithcatsAndNoCats.getSelected();
peopleWithcatsAndNoCats.getRejected();
```

# Reuse Code

## Stream API

```
Stream<Integer> stream = Lists.mutable.with(1, 2, 3).stream();
Assert.assertEquals(1, stream.mapToInt(i -> i).min().getAsInt());
Assert.assertEquals(3, stream.mapToInt(i -> i).max().getAsInt());
```

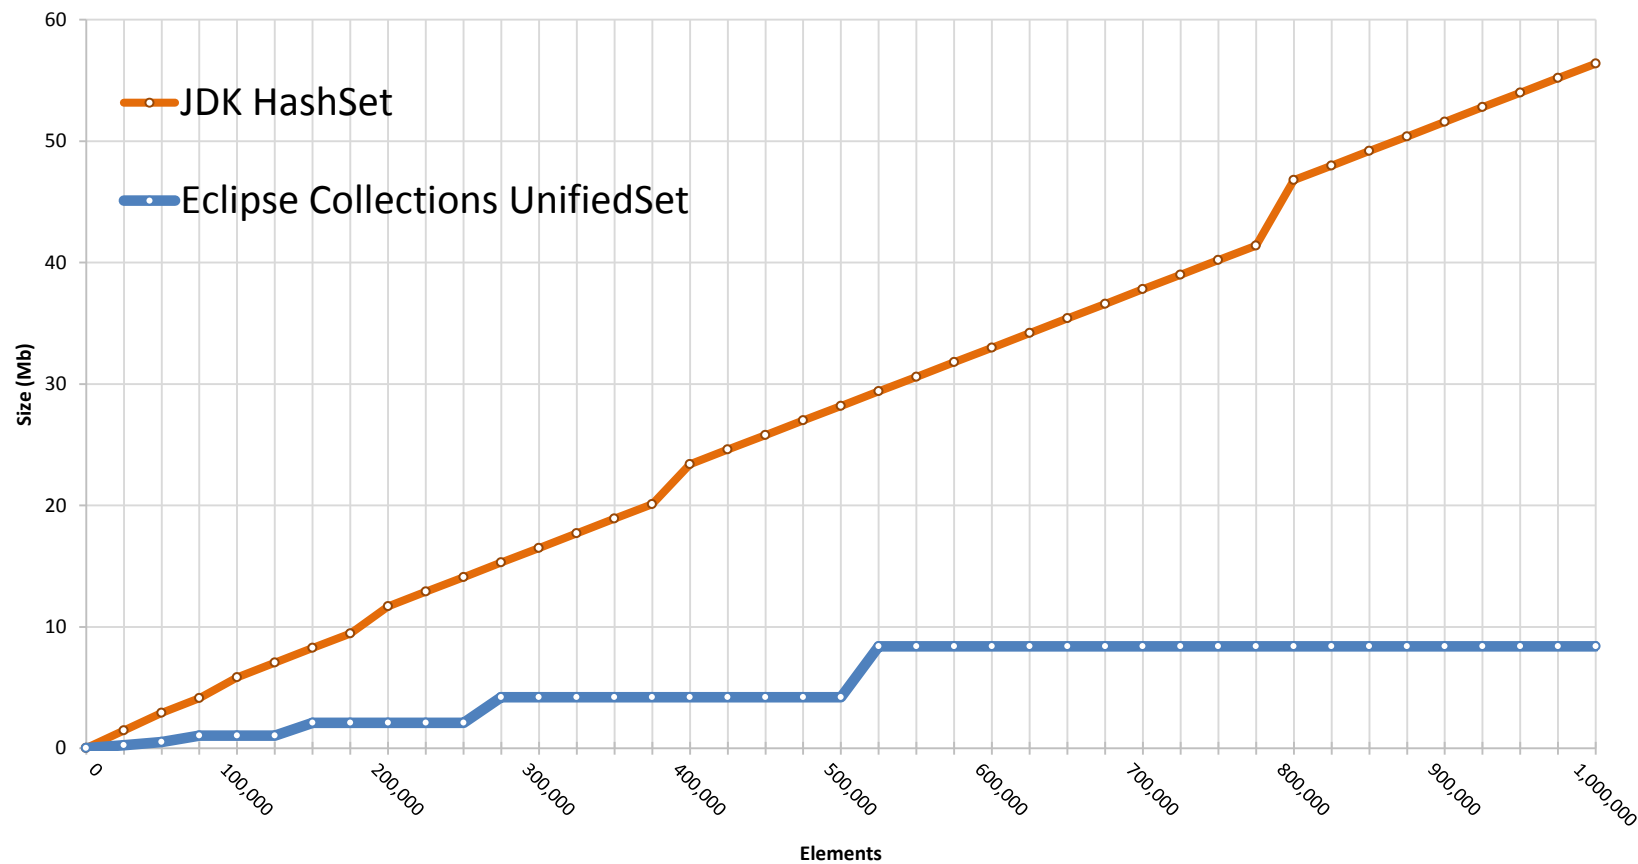
// throws

```
java.lang.IllegalStateException: stream has already been operated upon or closed
    at java.util.stream.AbstractPipeline.<init>(AbstractPipeline.java:203)
    at java.util.stream.IntPipeline.<init>(IntPipeline.java:91)
    at java.util.stream.IntPipeline$StatelessOp.<init>(IntPipeline.java:592)
    at java.util.stream.ReferencePipeline$4.<init>(ReferencePipeline.java:204)
    at java.util.stream.ReferencePipeline.mapToInt(ReferencePipeline.java:203)
```

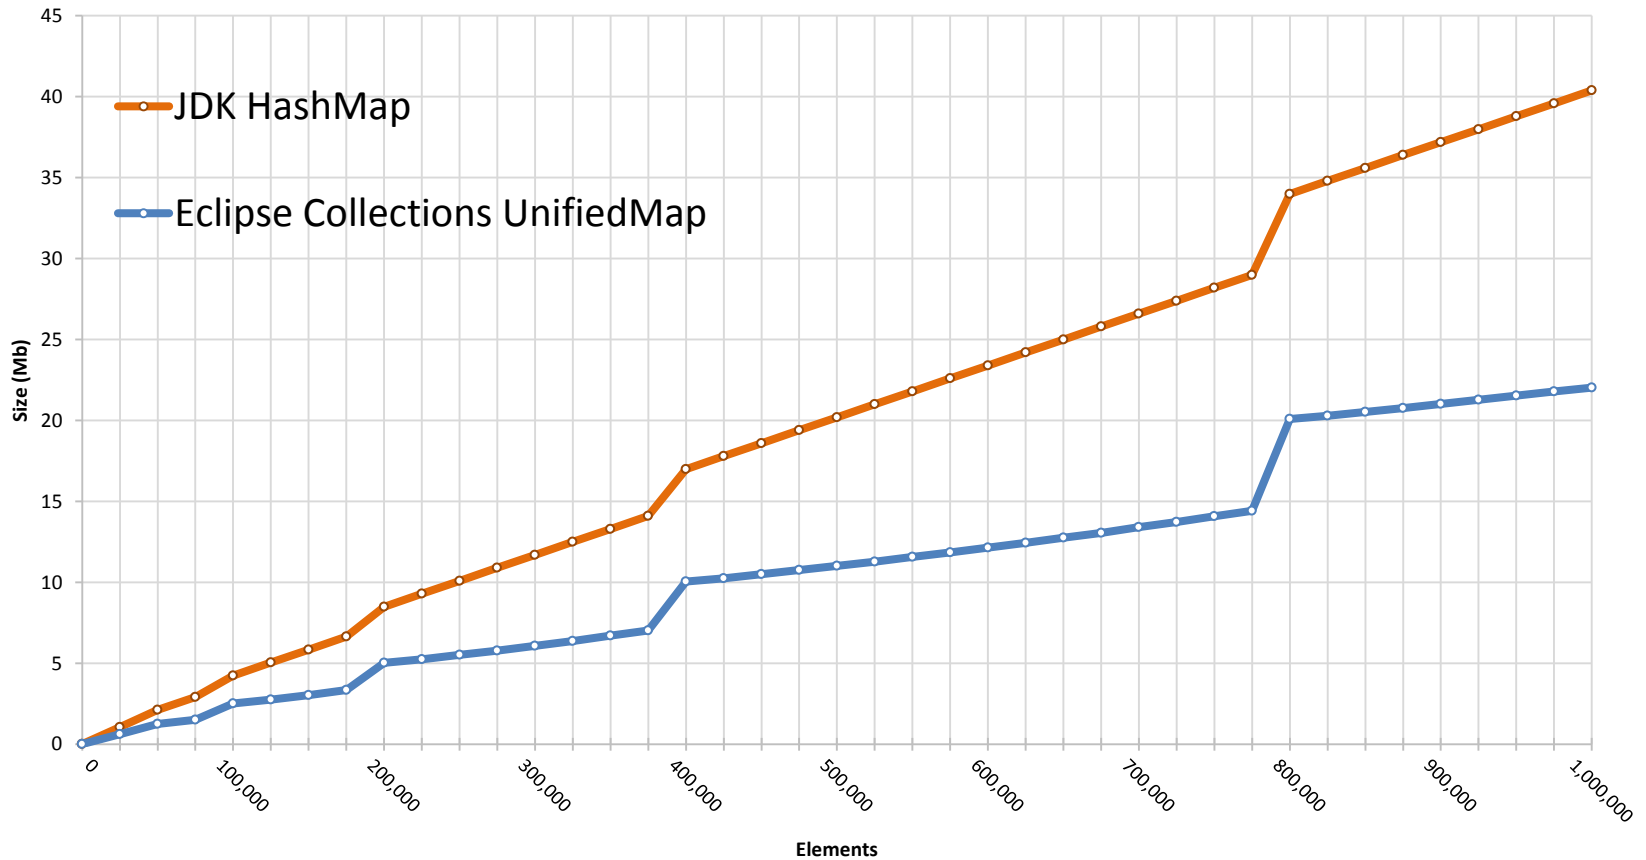
## Eclipse Collections

```
LazyIterable<Integer> lazy = Lists.mutable.with(1, 2, 3).asLazy();
Assert.assertEquals(1, lazy.collectInt(i -> i).min());
Assert.assertEquals(3, lazy.collectInt(i -> i).max());
```

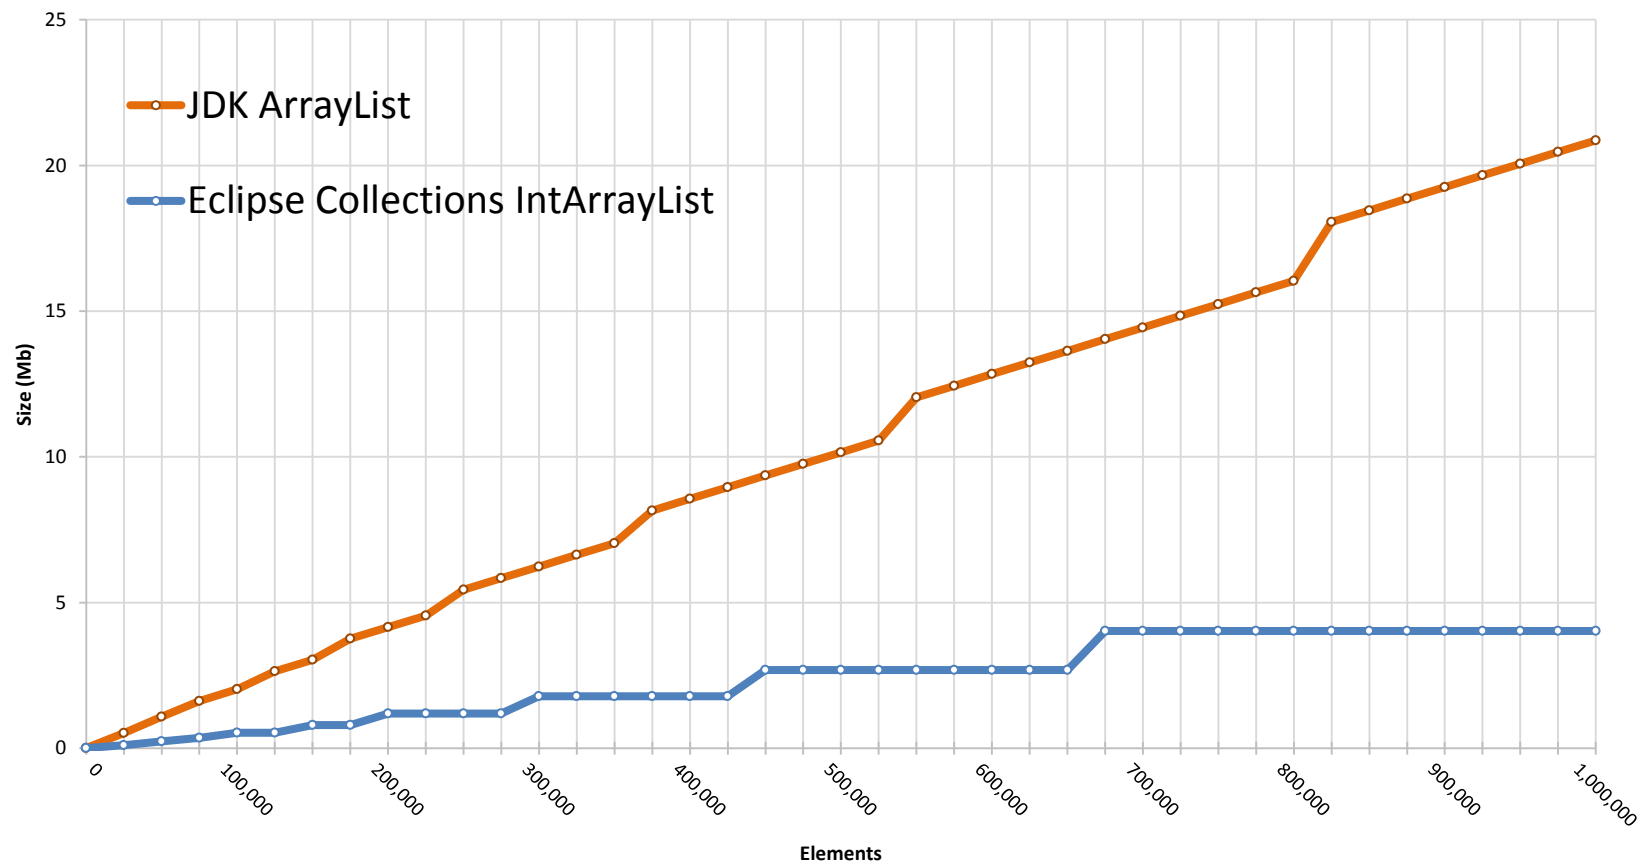
# Memory Efficient (HashSet vs UnifiedSet)



# Memory Efficient (HashMap vs UnifiedMap)



# Memory Efficient (primitive collections)



# Concise Exception Handling

Java  
Collections

```
Appendable appendable = new StringBuilder();
List<String> jdkList = Arrays.asList("one");
// jdkList.forEach(appendable::append);
// Compile Error: incompatible thrown types java.io.IOException in method reference
jdkList.forEach(each -> {
    try
    {
        appendable.append(each);
    }
    catch (IOException e)
    {
        throw new RuntimeException(e);
    }
});
Assert.assertEquals("one", appendable.toString());
```

Eclipse  
Collections

```
ImmutableList<String> eclipseList = Lists.immutable.with("two");
eclipseList.each(Procedures.throwing(appendable::append));
Assert.assertEquals("onetwo", appendable.toString());
```

# Appendix



# OpenJDK Enhanced Arrays.sort() JMH Results

## Test Sample

An array of int with size 10,000,000, 32 pair flips, a run of zeroes in the middle, and 32 pair flips at the end

Benchmark	(listType)	Mode	Cnt	Score	Error	Units
Sort(JDK8)	pairFlipZeroPairFlip	thrpt	10	4.886	± 0.031	ops/s
Sort(OpenJDK9)	pairFlipZeroPairFlip	thrpt	10	14.793	± 0.217	ops/s

Name	Description
Score	The mean for the throughput benchmarking mode
Error	Estimate the "noisyness" of the measurements taken by the microbenchmark
Thrpt	Mode.Throughput. Calculate number of operations in a time unit.
Ops/s	Operations per second

# Eclipse Collections: Patterns

## Eager iteration pattern

```
Interval list = Interval.fromTo(0, 100);  
list.detect(each -> each > 50);  
list.select(each -> each > 50);  
list.reject(each -> each > 50);  
list.anySatisfy(each -> each > 50);  
list.allSatisfy(each -> each > 50);  
list.collect(Object::toString);  
list.injectInto(3, Integer::sum);
```

## Lazy iteration pattern

```
list.asLazy()  
  .select(each -> each > 95)  
  .collect(Object::toString)  
  .makeString("[", ",", "]);  
  
=> [96,97,98,99,100]
```

# Parallel-lazy Performance:

## Java 8 vs Scala vs Eclipse Collections



<http://www.infoq.com/presentations/java-streams-scala-parallel-collections>

The session page got the top page view after QCon New York

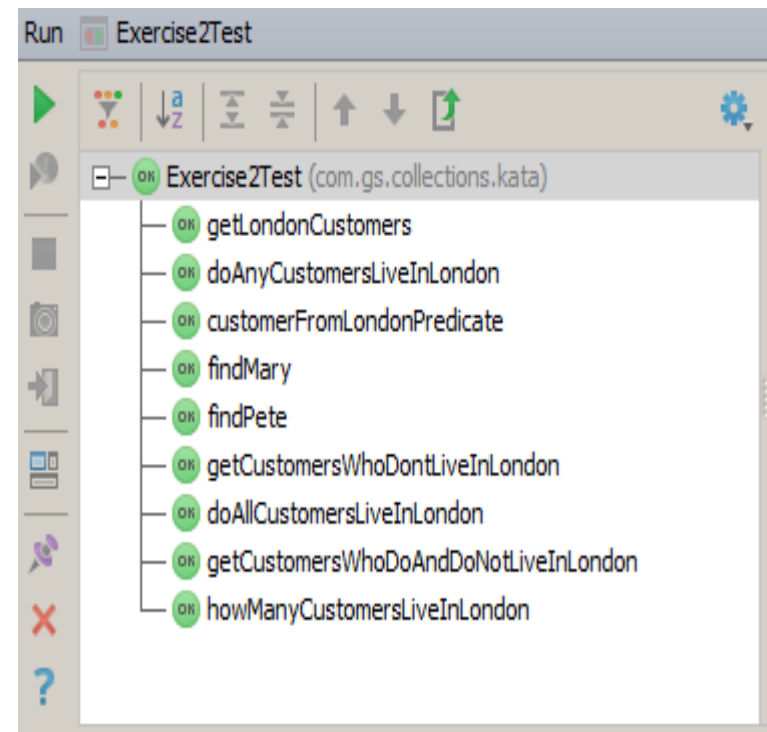
Here are the top 10 presentations from QCon New York 2014, based on page views:

- [Parallel-lazy Performance: Java 8 vs Scala vs GS Collections - Craig Motlin](#)
- [Evolving Java - Brian Goetz](#)
- [Caml Trading - Experiences with OCaml on Wall Street - Yaron Minsky](#)
- [End to End Reactive Programming at Netflix - Jafar Husain, Matthew Podwysocki](#)
- [Spring 4 on Java 8 - Juergen Hoeller](#)
- [10 Reasons Why Developers Hate Your API - John Musser](#)
- [What's the Best Way to Improve Software Architectures? - Eric Evans, Michael Feathers, Duncan DeVore, Leo Gorodinski](#)
- [How Facebook Scales Big Data Systems - Jeff Johnson](#)
- [A Day in the Life of a Functional Data Scientist - Richard Minerich](#)
- [Whither Web programming - Gilad Bracha](#)

# Training - Eclipse Collections Kata

<https://github.com/eclipse/eclipse-collections-kata>

- A training material by passing failed unit test one by one in TDD way
- It is also used in Goldman Sachs trainings
- Useful to learn the usage patterns of Eclipse Collections



# Scala Collections Performance

<http://www.goldmansachs.com/gs-collections/presentations/2015-03-17%20Scala%20Days%202015%20-%20Scala%20Collections%20Performance.pptx>

- Scala Days San Francisco 2015
- Performance analysis for Scala collections and other implementations
- Suggestions for Scala collection implementation to improve its performance, based on the knowledge we acquire from Eclipse Collections experience

# Eclipse Collections accepts contributions

**How To Contribute:** <https://github.com/eclipse/eclipse-collections/blob/master/CONTRIBUTING.md>

- Summary: Sign Eclipse Foundation CLA and submit pull-requests
- Developer mailing list:  
<https://dev.eclipse.org/mailman/listinfo/collections-dev>
- Reporting Issues: <https://github.com/eclipse/eclipse-collections/issues>

# Resources

- Eclipse Collections on GitHub  
<http://www.infoq.com/cn/news/2016/01/GS-Collections-Eclipse-Foundn>
- Eclipse Collections Memory Benchmark  
[http://www.goldmansachs.com/gs-collections/presentations/GSC\\_Memory\\_Tests.pdf](http://www.goldmansachs.com/gs-collections/presentations/GSC_Memory_Tests.pdf)
- Learn-by-Example: Eclipse Collections  
<http://www.infoq.com/cn/articles/gs-collections-examples-tutorial-part1>  
<http://www.infoq.com/cn/articles/gs-collections-examples-tutorial-part2>  
Eclipse Collections Kata <https://github.com/eclipse/eclipse-collections-kata>



# We BUILD

Learn more at [GS.com/Engineering](https://www.gs.com/engineering)

© 2016 Goldman Sachs. This presentation should not be relied upon or considered investment advice. Goldman Sachs does not warrant or guarantee to anyone the accuracy, completeness or efficacy of this presentation, and recipients should not rely on it except at their own risk. This presentation may not be forwarded or disclosed without Goldman Sachs' consent.

© 2016 Goldman Sachs. 本文章不得被依赖或被视为投资建议。高盛公司不担保、不保证本文章的精确、完整或效用。接收者不应依赖本文章，除非在自担风险的范围内。在未获得高盛公司同意的情形下，本文章不得被转发、披露。